

Stack Exchange Tag Prediction through Keyword Extraction

Gio Borje
University of California, Irvine
gborje@uci.edu

Paul Kang
University of California, Irvine
pdkang@uci.edu

ABSTRACT

Stack Exchange is a set some of the most renown, community-driven question and answer sites. The annotation of questions with tags enables users to find and respond to questions of interest immediately. However, the annotation of questions with tags is tedious and error-prone because users may not be aware of the best tags to categorize their question. By automating the process of annotating questions with tags, the community is relieved of their crowdsourcing work and enables them to focus on the important aspect of asking and answering questions. We survey several approach to keyword extraction and tag suggestion to Stack Exchange posts using a Bag-of-Words (BOW) model and show that Support Vector Machines have the best accuracy.

1. INTRODUCTION

Kaggle is a community of data scientists focused on solving complex data-science problems by providing company-hosted, public, data science competitions. Competitions proceed as follows. A company begins hosting a competition by providing a public data set as well as an evaluation metric. Participants register for the competition, download the data set, manipulate the data, and creates an appropriate model. Then, the participants download the test data set and make predictions on the test data using the trained models. The predictions for the test data are uploaded into the competition page, scored by the specified evaluation metric, and ranked accordingly on a public leaderboard.

This challenge is proposed through Kaggle: given only the title and body of a question, predict the question's tags. The training data set provided comes from Stack Exchange sites which mixes both technical and nontechnical questions. We do not solve this challenge. Instead, we survey a potential set of classifiers to see which classifier performs uniformly better than all others given the data set. Specifically, we evaluate a Bernoulli naïve Bayes classifier, Linear Support Vector Machine, Random Forest classifier, and a Gradient Boosting Machine.

The evaluation metric for this competition uses the Mean F_1 -Score which is commonly used to measure classification accuracy through *precision* and *recall*. Precision

can be defined as the ratio of true positive classifications over the true and false positive classifications. The recall can be defined as the ratio of the true positive classifications over the true positives and false negative classifications. The Mean F_1 -Score formula is then

$$F_1 = 2 \frac{pr}{p + r}$$

where p and r are precision and recall respectively. The Mean F_1 -Score is maximized through good precision and recall.

1.1 Data Set Analysis

First, we begin with a preliminary analysis of the Stack-Overflow data set. The training data contains three attributes: question ID, body and tags. The testing data contains two attributes: question ID and body while the tags are to be predicted. The body of each question contains raw HTML which also separates the title and body of each question.

Since the challenge fixes the number of tags for a question to be in between 1 and 5, we can utilize the distribution of the training data to suggest a particular number of tags. The distribution of the number of tags per questions can be seen in Table 1.1. It is necessary to realize that the

Number of Tags	Questions	Percentage
1	492518	13.76
2	951594	26.65
3	1023028	28.65
4	685565	19.17
5	420319	11.77

Table 1: Distribution of the number of tags per question

distribution is similar to a normal distribution which can be seen in Figure 1.1. A useful descriptive statistic is the average number of tags per question which has been observed to be 2.86.

1.2 Outline

This paper will proceed as follows. Section 2 will cover the technical details of our methods including the tools involved, the model selection process and a brief description

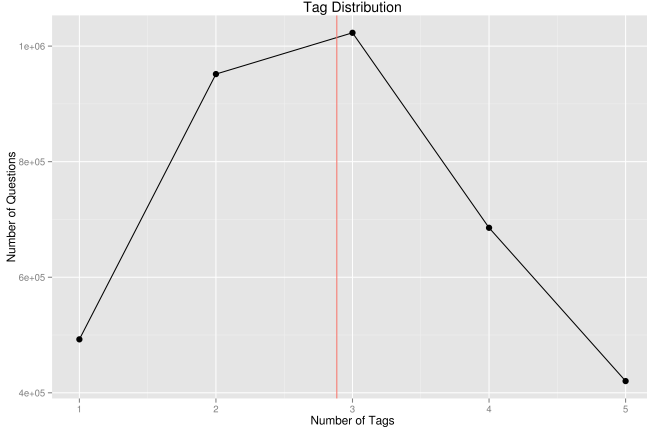


Figure 1: Distribution of the number of tags per question

of each model under evaluation. Section 3 will show our results of evaluated mean F_1 scores for each classifier. Section 4 will interpret our results including speculations of why Random Forests perform poorly with Latent Semantic Analysis. Finally, Section 5 will conclude with planned future work.

2. METHODS

2.1 Tools

The solution is programmed using Python. The implementations of the models, cross-validation and metrics are provided in the scikit-learn library which is coupled with scikit and numpy.

2.2 Problem Decomposition

We begin by decomposing the problem into a set of binary classifications tasks: for each tag, generate a classifier that predicts whether a given question should have the tag. Each tag must therefore have its own training set. We design the training set for each tag such that there are exactly 60 questions with the tag and 1500 questions without the tag. It is then easy to evaluate which classifier performs best for a specific tag.

2.3 Model Selection

In order to select the best model for each tag, several models are assessed using cross-validation such that the model with the best mean score is used as the classifier for the tag. In particular, we use stratified k -Fold cross-validation that partitions the training data to k folds such that the expected value of the predictions are approximately similar across each fold. We select k to be three folds since 60 and 1500 are equally divisible by three.

Each model assessed uses grid search for hyperparameter optimization. That is, the parameters of each model are

tuned exhaustively with bounds. Once all parameter combinations have been assessed, the parameters with the best mean cross-validation score are used in their model.

2.4 Preprocessing

Since we are using a Bag-of-Words (BOW) approach, we must vectorize the data set. That is, for each question, we use a count vectorizer which generates a histogram for each of the words in a question. The vectorized form is then used as the feature set for training models and making predictions.

A vectorizer's transformation of the data proceeds as follows. First, given the corpus of questions, it is tokenized and each token is used as a feature. Then, a histogram of the tokens are computed per question and stored as rows of a matrix. Hence, this matrix should have dimensions $N \times M$ where N is the number of questions and M is the number of tokens extracted. We call this matrix, the *question-token matrix*. We can visualize this matrix as seen on Figure 2.

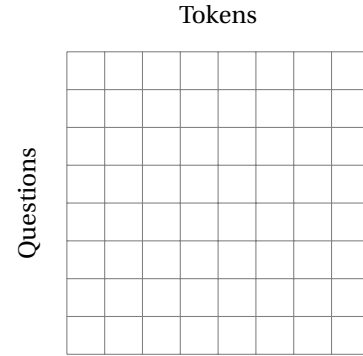


Figure 2: Question-Token Matrix

Once the question-token matrix has been computed, an additional preprocessing step follows: Latent Semantic Analysis. That is, we apply Singular Value Decomposition (SVD) on the matrix to reduce the dimensionality. We specify the reduced dimensionality to be 512 for rapid prototyping.

2.5 Bernoulli Naïve Bayes

The naïve Bayes approach applies Bayes' theorem in order to classify. That is, predicting existence of a tag in a question can be reformulated as follows:

$$p(y|\mathbf{x}) = \frac{p(y)p(\mathbf{x}|y)}{p(\mathbf{x})}$$

where y is a boolean value representing whether or not the tag should exist and \mathbf{x} is the feature vector of the question. In order to compute the classifying value, y , we reformulate the problem as an optimization problem:

$$\hat{y} = \underset{y}{\operatorname{argmax}} p(y) \prod_{i=1}^n p(x_i|y)$$

where the \hat{y} is the classification made.

The Bernoulli variant of naïve Bayes assumes that the data is distributed according to multivariate Bernoulli distributions. This particular implementation is best applied to data sets with a small vocabulary (under 1000 features).[4] The decision rule is based on the following formula:

$$p(x_i|y) = p(i|y)x_i \times (1 - p(i|y))(1 - x_i)$$

In our solution, we hyperoptimize $\alpha \in \{0, 1, 2\}$ which is an additive smoothing parameter.

2.6 Random Forests

The Random Forest approach is an ensemble method that uses a set of weak learners to form a strong learner.[1] The weak learners are frequently decision trees which are grouped together into a forest. The number of decision trees generated is parameterized by T . Each split of a decision tree is computed starting with a random subset of the features of size m and selecting the best of those features.

The classification process in Random Forests is as follows. First, the feature vector is input into each of the trees of the forest. Then, the classification of the classifier should be the class that is the voting majority of the trees in the forest.

In our solution, we hyperoptimize $m \in \{\sqrt{N}, \log_2 N\}$ where N is the number of features and we also hyperoptimize T .

2.7 Linear Support Vector Classifier

The idea of a Support Vector Machine is to define a separating hyperplane which distinguishes a set of classes, which, in this case, is a boolean value of whether or not a tag should be accepted. There are, however, many feasible separating hyperplanes. Hence, we define the optimal hyperplane to be the hyperplane that maximizes the *margin*. The margin is defined to be the largest minimum distance of the hyperplane to each of the training samples.

We can compute the optimal hyperplane by reformulating the problem as an optimization problem. The primal form of the optimization problem can be described as follows:

$$\operatorname{argmin}_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$$

where \mathbf{w} and b define the optimal hyperplane.

Since, in our data set, it is rarely the case that the data is linearly separable, the optimal hyperplane can be redefined using the *Soft Margin* method which

$$\operatorname{argmin}_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \xi_i$$

introduces non-negative slack variables, ξ_i , and a penalty parameter, C . [2] In our solution, we hyperoptimize $C \in \{10^i | 0 \leq i \leq 9\}$ using grid search.

2.8 Gradient Boosting Machines

Gradient Boosting Machines (GBM) is an ensemble method on weak classifier that builds its model in a stage-wise fashion where each stage minimizes a specified loss function using the Gradient Descent method.[3] Each weak classifier is described as a function $h(\mathbf{X}; \mathbf{p}_i)$ where \mathbf{X} is the input and \mathbf{p}_i is its parameter set. A loss function, $L(\mathbf{y}, \mathbf{y}')$ is then specified on \mathbf{y} and \mathbf{y}' which are the expected classification and predicted classification respectively. The predicted classification is computed as a linear combination of the weak learners:

$$\mathbf{y}' = \sum_{i=1}^K \beta_i h(\mathbf{X}; \mathbf{p}_i)$$

where K is the number of weak learners and β_i is a weight parameter for a specified weak learned, i .

Hence, the model is described by β and \mathbf{P} which are the weak learner weights and parameter sets respectively. We can then formulate the problem as an optimization problem:

$$\operatorname{argmin}_{\beta, \mathbf{P}} L(\mathbf{y}, \sum_{i=1}^K \beta_i h(\mathbf{X}; \mathbf{p}_i))$$

In our solution, we hyperoptimize $K \in \{100, 200, 300\}$, $\max_features \in \{\sqrt{M}, \log_2 M\}$ and $\min_samples_split \in \{1, 2, 3\}$.

3. RESULTS

The Mean F_1 -Score for the each of the models on the specified subset of tags are shown on Table 2 and Table 3 which represent scores for models using count vectorized features and Latent Semantic Analysis (LSA) respectively. With LSA, Bernoulli naïve Bayes and Gradient Boosting Machines perform significantly better; however, Random Forests perform significantly worse. The loss of accuracy in SVMs with LSA is negligible. Generally, SVMs yield the best results in both feature sets; however, with LSA, the naïve Bayes approach sometimes generates a better classifier. A graph of the results of classifier scores using LSA is shown in Figure 2.8.

We can rank classifiers according to their mean F_1 -score as seen in Figure 2.8. In order, the best classifiers are SVMs, Bernoulli naïve bayes, Gradient Boosting Machines and Random Forests in that order.

4. DISCUSSION

Overall, the classifiers had the best accuracy for specific and uncommon tags such as "sqlalchemy." Conversely, the classifiers had the worst accuracy for popular tags such as "c", "c++" and "python". It is interesting to note that "haskell" was accurately classified often. We believe that that popular tags are hard to classify because they can easily be associated with unrelated terms; in other words, there is a significant amount of noise associated with popular tags.

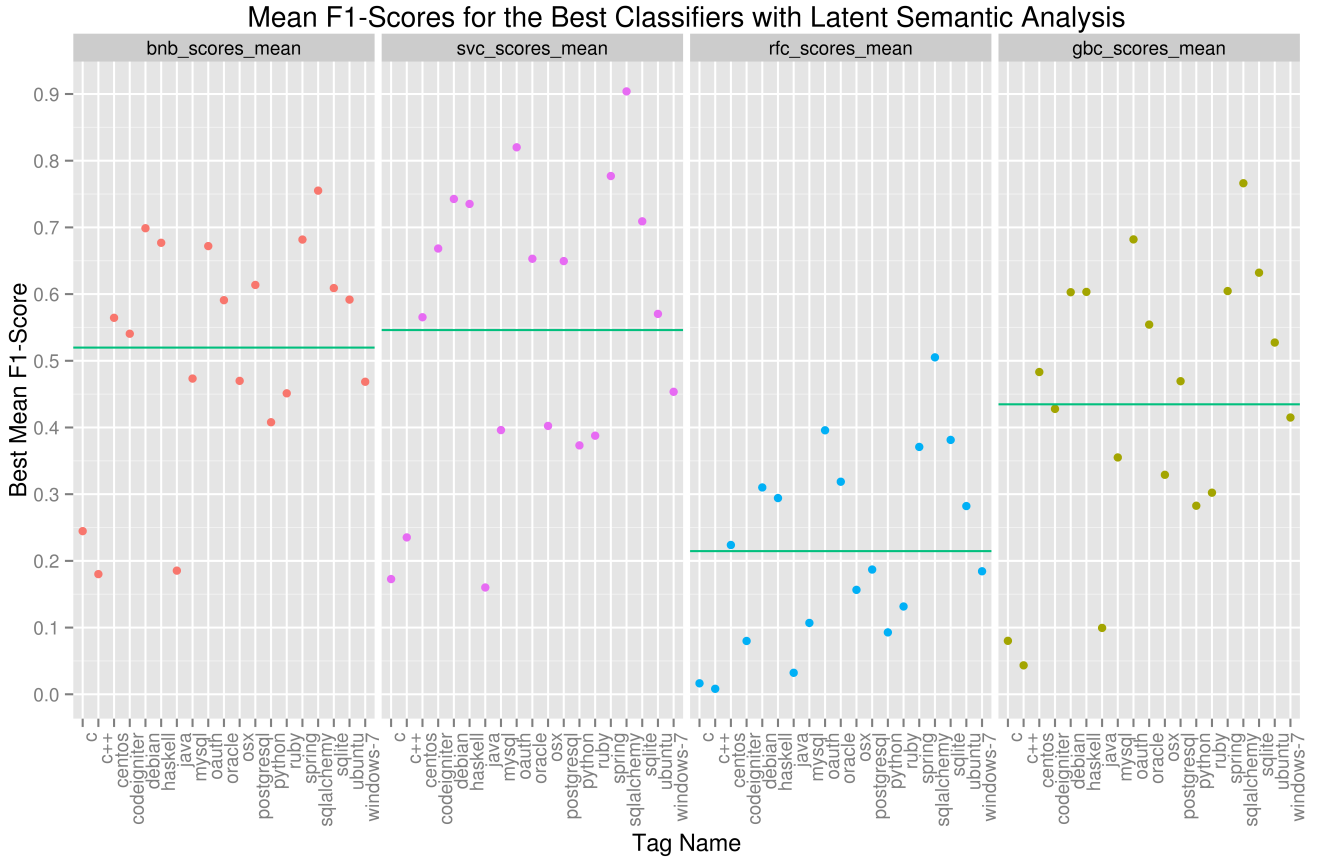


Figure 3: Mean F_1 -Score decomposition by Classifier

We believe that the data has a significant amount of high-dimensional noise which can be observed through the poor scores of nonlinear classifiers. That is, it may be the case that the nonlinear classifiers are overfitting due to the significant variance in the methods. The explained variance also supports our hypothesis that popular tags are polluted with noisy features. Therefore, on the next training phases, it is logical to hyperoptimize over parameters that reduce variance in the classifiers.

The introduction of Latent Semantic Analysis (LSA) made a significant impact on the majority of the classifiers. We suspect the loss of accuracy in Random Forests using LSA is due to the fact that Random Forests already performs a type of dimensionality reduction in its weak learners by randomly selecting subsets of the designated feature set. Hence, reducing dimensionality with LSA before fitting Random Forests yields redundant dimensionality reduction

5. CONCLUSION

Although the overall accuracy is low, we have established a baseline metric for evaluating further methods to

be used in this competition. There is much future work to be done such as using Parts-of-Speech as a feature set which Latent Semantic Analysis does not take into consideration. Another potential feature for programming languages in specific is to design specific language classifiers using the known publicly available syntax and grammar specifications of each language.

In order to solve the bias-variance dilemma, we may consider an ensemble method which produces a classification based on bagging or boosting the given models. The running time of learning strong models may, however, be a concern. Furthermore, Stochastic Gradient Descent Boosting is a natural consideration to tackle the issue of variance.

6. ACKNOWLEDGEMENTS

We would like to thank the Stack Exchange communities for making available their data set.

7. REFERENCES

- [1] Leo Breiman. Random forests. *Machine learning*, 45(1):5–32, 2001.

- [2] Corinna Cortes and Vladimir Vapnik. Support-vector networks. *Machine learning*, 20(3):273–297, 1995.
- [3] Jerome H Friedman. Greedy function approximation: a gradient boosting machine. *Annals of Statistics*, pages 1189–1232, 2001.
- [4] Andrew McCallum, Kamal Nigam, et al. A comparison of event models for naive bayes text classification. In *AAAI-98 workshop on learning for text categorization*, volume 752, pages 41–48. Citeseer, 1998.

Tag	Bernoulli Naïve Bayes	Linear Support Vector Classifier	Random Forests	Gradient Boosting
codeigniter	0.1764	0.7202	0.5176	0.6122
spring	0.2957	0.7268	0.4874	0.6456
sqlalchemy	0.3516	0.8959	0.7742	0.8543
oauth	0.4117	0.8238	0.4157	0.6899
mysql	0.1256	0.3877	0.0240	0.2006
oracle	0.1817	0.6398	0.2893	0.5034
postgresql	0.1957	0.6299	0.5194	0.3807
sqlite	0.1731	0.7293	0.4919	0.5642
ubuntu	0.2694	0.5701	0.2040	0.4390
debian	0.3079	0.7239	0.5674	0.7042
centos	0.2353	0.5878	0.1641	0.3173
osx	0.1283	0.4061	0.1635	0.3469
windows-7	0.2498	0.4527	0.0465	0.2070
python	0.0921	0.3181	0.0951	0.1563
java	0.0727	0.1725	0.0081	0.0662
c++	0.1291	0.2110	0.0082	0.0830
c	0.1314	0.2962	0.0313	0.1168
ruby	0.0978	0.3972	0.1328	0.2048
haskell	0.2524	0.7708	0.7378	0.6077

Table 2: Mean F_1 -Scores of Models with Count Vectorized Feature Vectors

Tag	Bernoulli Naïve Bayes	Linear Support Vector Classifier	Random Forests	Gradient Boosting
codeigniter	0.5405	0.6683	0.0799	0.4278
spring	0.6817	0.7770	0.3708	0.6045
sqlalchemy	0.7551	0.9038	0.5051	0.7662
oauth	0.6719	0.8200	0.3958	0.6820
mysql	0.4733	0.3961	0.1069	0.3550
oracle	0.5908	0.6530	0.3187	0.5542
postgresql	0.6137	0.6495	0.1870	0.4694
sqlite	0.6090	0.7090	0.3813	0.6321
ubuntu	0.5917	0.5702	0.2821	0.5273
debian	0.6987	0.7426	0.3100	0.6028
centos	0.5644	0.5653	0.2238	0.4831
osx	0.4699	0.4023	0.1565	0.3290
windows-7	0.4685	0.4535	0.1844	0.4149
python	0.4078	0.3731	0.0927	0.2826
java	0.1853	0.1601	0.0322	0.0995
c++	0.1801	0.2352	0.0082	0.0434
c	0.2446	0.1727	0.0164	0.0801
ruby	0.4512	0.3877	0.1316	0.3022
haskell	0.6769	0.7351	0.2941	0.6033

Table 3: Mean F_1 -Scores of Models with Latent Semantic Analysis